

Introduction: (from the group report)

Group 3 has decided to improve on a body of existing code that classifies a series of images, known as TrashNet. Our task is to improve model efficacy, facilitating better garbage sorting recognition. The ultimate goal is to improve the pre-trained models in our baseline code via transfer learning, select more modern models, then build an application which showcases our ability to correctly classify trash.

Detailed in this whitepaper will be a description of the data, background on the networks, our setup for experimentation, results, and conclusions.

Individual work entry(what I did and results):

→ (I am so sorry about the formatting but this was how I document my part when I was doing research/testing)

→ published all versions (mentioned below) into github JupiK9 branch (should be merged together now)

Data storage

Originally want to store it in S3 bucket using AWS. Went to create the bucket, only to find out that our deep learning class doesn't give permission to create S3 (also don't want to spend money) so ended up handing down to Stephen using GCP.

- Result: the S3 bucket not created successfully, because no permission guaranteed from the DL class- AWS account. Don't want to spend money
- Steps of setting up S3 bucket:

SafariFileEditViewHistoryBookmarksWindowHelp

沒必要的修辭手法堵住我的嘴

Getting started...Deep-Learning...Deep-Learning...My AppsAccounts | AWS...us-east-1.console.aws.amazon.com

Search[Option+S]

United States (N. Virginia)CCAS-DATS-DL-Student/G28823536@gsu.edu

Amazon S3> Buckets> Create bucket

Create bucketinfo

Buckets are containers for data stored in S3.

General configuration

AWS Region
US East (N. Virginia) us-east-1

Bucket typeinfo

☒ General purpose

Recommended for most use cases and access patterns. General purpose buckets are the original S3 bucket type. They allow a mix of storage classes that redundantly store objects across multiple Availability Zones.

☐ Directory

Recommended for low-latency use cases. These buckets use only the S3 Express One Zone storage class, which provides faster processing of data within a single Availability Zone.

Bucket nameinfo

sp25di

Bucket names must be 3 to 63 characters and unique within the global namespace. Bucket names must also begin and end with a letter or number. Valid characters are a-z, 0-9, periods (.), and hyphens (-). [Learn more](#)

Copy settings from existing bucket - optional

Only the bucket settings in the following configuration are copied.

☒ Choose bucket

Format: s3://bucket/profile

Object Ownershipinfo

Control ownership of objects written to this bucket from other AWS accounts and the use of access control lists (ACLs). Object ownership determines who can specify access to objects.

☒ ACLs disabled (recommended)

All objects in this bucket are owned by this account. Access to this bucket and its objects is specified using only policies.

☐ ACLs enabled

Objects in this bucket can be owned by other AWS accounts. Access to this bucket and its objects can be specified using ACLs.

Object Ownership

Bucket owner enforced

Block Public Access settings for this bucket

Public access is granted to buckets and objects through access control lists (ACLs), bucket policies, access point policies, or all. In order to ensure that public access to this bucket and its objects is blocked, turn on Block all public access. These settings apply only to this bucket and its access points. AWS recommends that you turn on Block all public access, but before applying any of these settings, ensure that your applications will work correctly without public access. If you require some level of public access to this bucket or objects within, you can customize the individual settings below to suit your specific storage use cases. [Learn more](#)

CloudShellFeedback

© 2025, Amazon Web Services, Inc. or its affiliates. PrivacyTermsCookie preferences

SafariFileEditViewHistoryBookmarksWindowHelp

沒必要的修辭手法堵住我的嘴

Getting started...Deep-Learning...Deep-Learning...My AppsAccounts | AWS...us-east-1.console.aws.amazon.com

Search[Option+S]

United States (N. Virginia)CCAS-DATS-DL-Student/G28823536@gsu.edu

Amazon S3> Buckets> Create bucket

Block Public Access settings for this bucket

Public access is granted to buckets and objects through access control lists (ACLs), bucket policies, access point policies, or all. In order to ensure that public access to this bucket and its objects is blocked, turn on Block all public access. These settings apply only to this bucket and its access points. AWS recommends that you turn on Block all public access, but before applying any of these settings, ensure that your applications will work correctly without public access. If you require some level of public access to this bucket or objects within, you can customize the individual settings below to suit your specific storage use cases. [Learn more](#)

☐ Block all public access

Turning this setting on is the same as turning on all four settings below. Each of the following settings are independent of one another.

☐ Block public access to buckets and objects granted through new access control lists (ACLs)

S3 will block public access permissions applied to newly added buckets or objects, and prevent the creation of new public access ACLs for existing buckets and objects. This setting doesn't change any existing permissions that allow public access to S3 resources using ACLs.

☐ Block public access to buckets and objects granted through any access control lists (ACLs)

S3 will ignore all ACLs that grant public access to buckets and objects.

☐ Block public access to buckets and objects granted through new public bucket or access point policies

S3 will block new bucket and access point policies that grant public access to buckets and objects. This setting doesn't change any existing policies that allow public access to S3 resources.

☐ Block public and cross-account access to buckets and objects through any public bucket or access point policies

S3 will ignore public and cross-account access for buckets or access points with policies that grant public access to buckets and objects.

⚠ Turning off block all public access might result in this bucket and the objects within becoming public

AWS recommends that you turn on block all public access, unless public access is required for specific and verified use cases such as static website hosting.

☒ I acknowledge that the current settings might result in this bucket and the objects within becoming public.

Bucket Versioning

Versioning is a means of keeping multiple variants of an object in the same bucket. You can use versioning to preserve, retrieve, and restore every version of every object stored in your Amazon S3 bucket. With versioning, you can easily recover from both unintended user actions and application failures. [Learn more](#)

Bucket Versioning

☒ Disable

☐ Enable

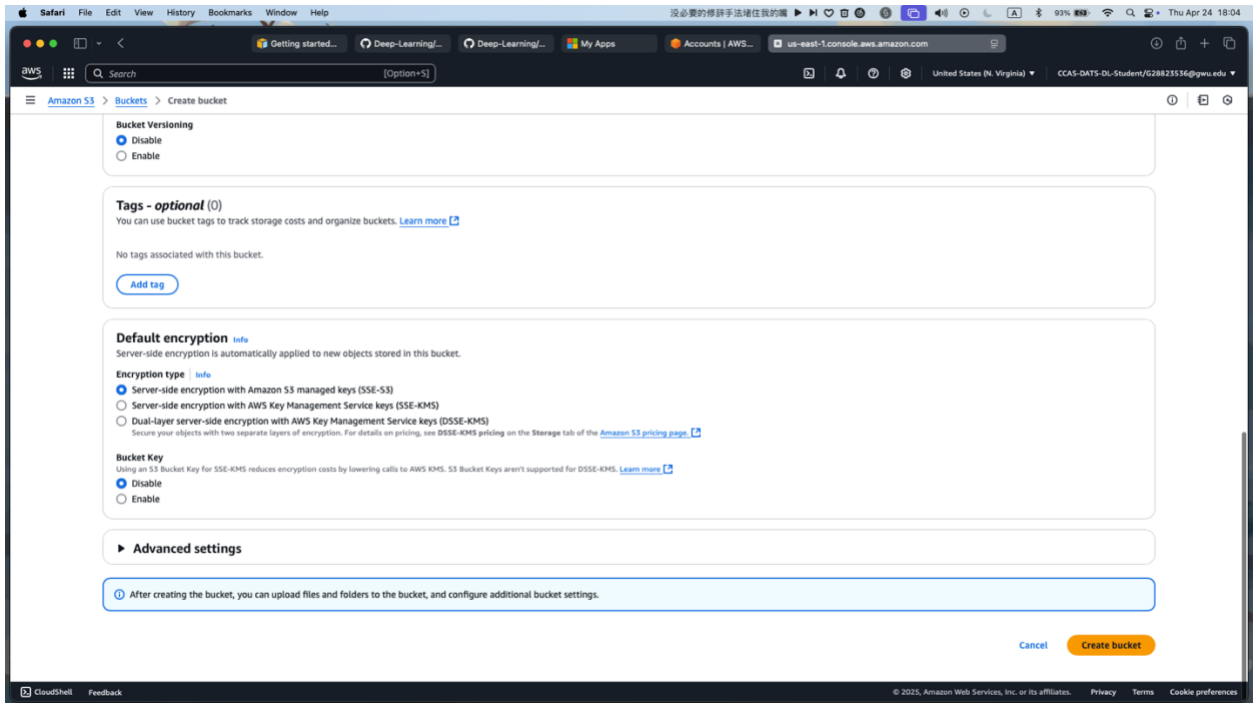
Tags - optional (0)

You can use bucket tags to track storage costs and organize buckets. [Learn more](#)

No tags associated with this bucket.

CloudShellFeedback

© 2025, Amazon Web Services, Inc. or its affiliates. PrivacyTermsCookie preferences



Improving baseline code -v2

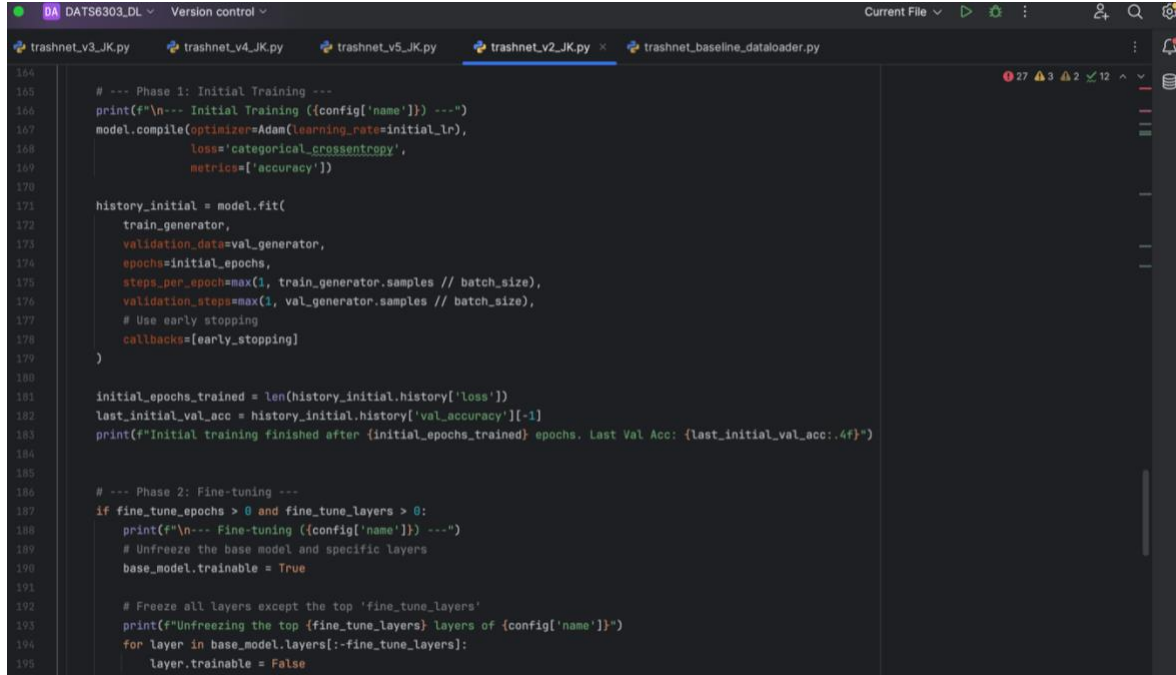
Experimental setups: including training strategy, hyperparameter control, regularization, data splitting, data augmentation and optimizer.

Training strategy:

- Two phase training:
 - Phase 1 – initial training: trains only the new classification head for `initial_epochs`. Quickly trains the new layers to adapt to the dataset based on general features
 - Phase 2 – fine-tuning: selectively freezes all layers except the top `fine_tune_layers`. Recompiled with low learning rate and trained for `additional_fine_tune_epochs`. → allows the model adjust pre-trained weights

in the later layers, becoming more complex and leads to significant accuracy improvements.

- Using low learning rate is crucial here to avoid catastrophically disrupting the valuable pre-trained weights



```
164
165 # --- Phase 1: Initial Training ---
166 print(f"\n--- Initial Training ({config['name']}) ---")
167 model.compile(optimizer=Adam(learning_rate=initial_lr),
168               loss='categorical_crossentropy',
169               metrics=['accuracy'])
170
171 history_initial = model.fit(
172     train_generator,
173     validation_data=val_generator,
174     epochs=initial_epochs,
175     steps_per_epoch=max(1, train_generator.samples // batch_size),
176     validation_steps=max(1, val_generator.samples // batch_size),
177     # Use early stopping
178     callbacks=[early_stopping]
179 )
180
181 initial_epochs_trained = len(history_initial.history['loss'])
182 last_initial_val_acc = history_initial.history['val_accuracy'][-1]
183 print(f"Initial training finished after {initial_epochs_trained} epochs. Last Val Acc: {last_initial_val_acc:.4f}")
184
185
186 # --- Phase 2: Fine-tuning ---
187 if fine_tune_epochs > 0 and fine_tune_layers > 0:
188     print(f"\n--- Fine-tuning ({config['name']}) ---")
189     # Unfreeze the base model and specific layers
190     base_model.trainable = True
191
192     # Freeze all layers except the top 'fine_tune_layers'
193     print(f"Unfreezing the top {fine_tune_layers} layers of {config['name']}")
194     for layer in base_model.layers[:-fine_tune_layers]:
195         layer.trainable = False
```

- hyperparameter control and flexibility
 - introduces more specific hyperparameters: initial_epochs, fine_tune_epochs, initial_lr, fine_tune_lr, fine_tune_layers.
 - To precisely control the two phase training process
 - Separate learning rates are crucial for fine-tuning
 - Specifying the number of layers to unfreeze the fine_tune_layers provides control over how much of the base model is adapted.
- Regularization
 - Added a dropout layer before the final output layer
 - Helps prevent overfitting problem
 - How: randomly set a fraction of neuron outputs to 0 during training, forcing the network to learn more robust features
 - Implements earlystopping during both training phases.
 - Monitors the validation loss or accuracy

- Stops training if it doesn't improve for a set number of epochs: `early_stopping_patience`
- Prevents overfitting by stopping before the model starts performing worse on unseen data,
- Save training time
- Restores the weights from the best performing epoch

Result (only with two model configs: ResNet101V2 and MobileNetV2)

```

# model.save(model_save_path)
# print(f"Saving history to {history_save_path}")
# np.save(history_save_path, history_final_history)

```

Terminal Local x Remote Python 3.12.3 (http://ubuntu@3.84.69.109:22/...) x + v

```

59/59 ----- 5s 87ms/step - accuracy: 0.6562 - loss: 0.8028 - val_accuracy: 0.6198 - val_loss: 0.8469
Epoch 3/10
59/59 ----- 20s 108ms/step - accuracy: 0.7935 - loss: 0.5983 - val_accuracy: 0.7533 - val_loss: 0.6738
Epoch 4/10
59/59 ----- 0s 2ms/step - accuracy: 0.8125 - loss: 0.6078 - val_accuracy: 0.8571 - val_loss: 0.5789
Epoch 5/10
59/59 ----- 20s 510ms/step - accuracy: 0.8485 - loss: 0.4717 - val_accuracy: 0.7747 - val_loss: 0.6229
Epoch 5: early stopping
Restoring model weights from the end of the best epoch: 1.
Initial training finished after 5 epochs. Last Val Acc: 0.7747

--- Fine-tuning (ResNet101V2) ---
Unfreezing the top 20 Layers of ResNet101V2
Epoch 6/20
59/59 ----- 45s 471ms/step - accuracy: 0.7381 - loss: 0.8809 - val_accuracy: 0.7336 - val_loss: 0.7238
Epoch 7/20
1/59 ----- 2s 47ms/step - accuracy: 0.6875 - loss: 0.68832025-04-26 23:17:14.857166: I tensorflow/core/framework/local_rendezvous.cc:484] Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence
[[[Node IteratorGetNext]]]]
59/59 ----- 2s 31ms/step - accuracy: 0.6875 - loss: 0.6883 - val_accuracy: 0.6198 - val_loss: 0.7117
Epoch 8/20
59/59 ----- 20s 397ms/step - accuracy: 0.7981 - loss: 0.6438 - val_accuracy: 0.7319 - val_loss: 0.7113
Epoch 9/20
59/59 ----- 0s 2ms/step - accuracy: 0.7812 - loss: 0.7316 - val_accuracy: 0.6667 - val_loss: 0.9076
Epoch 10/20
59/59 ----- 20s 397ms/step - accuracy: 0.8269 - loss: 0.5848 - val_accuracy: 0.7588 - val_loss: 0.6943
Epoch 10: early stopping
Restoring model weights from the end of the best epoch: 6.

--- Evaluating on Test Set (ResNet101V2) ---
15/15 ----- 1s 44ms/step - accuracy: 0.7678 - loss: 0.5744

Model Comparison Results (with Fine-tuning):

```

	Model	Test Accuracy	Test Loss	Final Train Accuracy	Final Val Accuracy	Initial Epochs	FineTune Epochs
1	ResNet101V2	0.7667	0.6857	0.8119	0.7588	5	8
0	MobileNetV2	0.7271	0.7048	0.7492	0.6998	7	-2

ubuntu@ip-18-1-3-133:~/Jupiter_d1/Final_project

1. V3

- slightly changed the test split from the baseline code to 20%.
- dataset is not big enough and we have to use all we have, the original baseline did a second ImageDataGenerator, in which it has its own train-val 75/25.
- leaving it as a potential data leakage problem, due to the independent application of val_split by each generator on the full dataset,
- test samples included in either train or validation set of the second generator, as well as leading overlap and miscounts.

→ the day we have proper data storage and data loader

only add back all the models

Result of having all model config trained, tested and validated.

Model Comparison Results (with Fine-tuning):									
	Model	Test Accuracy	Test Loss	Final Train Accuracy	Final Val Accuracy	Initial Epochs	FineTune Epochs		
2	ResNet152V2	0.7896	0.6110	0.7974	0.7516	5	0		
1	ResNet101V2	0.7688	0.6442	0.8151	0.7566	5	2		
0	MobileNetV2	0.7292	0.7354	0.7401	0.6908	7	-2		
3	MobileNet	0.6729	0.8270	0.7176	0.6595	5	0		
5	MobileNetV3Large	0.3583	1.4839	0.3248	0.3257	5	0		
4	MobileNetV3Small	0.2542	1.6552	0.2219	0.2549	5	0		

V4:

- optimizer: to AdamW, added weight_decay hyperparameter
 - o decouples weight decay from the adaptive learning rate adjustments
 - o more effective way to apply regularization, can lead to better model generalization and final performance

For v4.py, some of the improvements are:

- Increase epoch slightly, for initial, fine_tune;
- Slightly decrease fine_tune LR

- Unfreeze more layers
- Increase patience for callback
- Weighted decay parameter for AdamW
-
- Add another dense layer and dropout to the head: (line 168)

```
# Add new dense layer:
x = GlobalAveragePooling2D(name="global_avg_pooling")(base_model.output)
x = Dense(256, activation='relu', name="dense_256")(x)
x = Dropout(dropout_rate, name="dropout_1")(x)
x = Dense(128, activation='relu', name="dense_128")(x)
x = Dropout(dropout_rate, name="dropout_2")(x)
outputs = Dense(6, activation='softmax', name="output_softmax")(x)
```

-
-
- In training loop – phase 2:
 - Freeze all layers except for the top fine_tune_layer
 - Use AdamW optimizer with fine-tune LR and weight decay

```
# Use AdamW optimizer
optimizer_finetune = AdamW(learning_rate=fine_tune_lr, weight_decay=weight_decay)
model.compile(optimizer=optimizer_finetune,
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

-
- Sort the result as best val accuracy before looking at the test accuracy
- Result::

```
Model Comparison Results (Improved Tuning):
Model Test Accuracy Test Loss Best Val Accuracy Final Train Accuracy Initial Epochs Run FineTune Epochs Run
0 MobileNetV2 0.7354 0.6738 0.9048 0.7460 15 -8
2 ResNet152V2 0.7625 0.7249 0.8095 0.6897 7 0
1 ResNet101V2 0.7417 0.7820 0.7812 0.6190 7 0
3 MobileNet 0.6500 0.8861 0.7023 0.5740 7 0
4 MobileNetV3Small 0.3417 1.6610 0.4286 0.2047 7 0
5 MobileNetV3Large 0.4021 1.5741 0.4013 0.2524 7 0
ubuntu@ip-10-1-3-133:~/Jupiter_d1/Final_project/DEEP_LEARNING_6303_GROUP3/Code/Trashnet/code$
```

V5

Stephen made a class weight parameter, fine tuning in alpha parameter. Because we never actually classified them, so they've been underrepresented. After adding the class weights by changing the alpha level, (that's when I started to play around with the data)

By increasing alpha value from 0.2 to 0.5: the result increases:

Model Comparison Results (with Fine-tuning):									
	Model	Test Accuracy	Test Loss	Final Train Accuracy	Final Val Accuracy	Initial Epochs	FineTune Epochs		
0	EfficientNetV2S	0.8458	0.4713	0.8934	0.8248	18	-5		
2	ResNet101V2	0.7750	0.6347	0.7958	0.7385	5	0		
3	ResNet152V2	0.7688	0.6384	0.7958	0.7253	5	0		
1	MobileNetV2	0.7229	0.7663	0.7224	0.6957	5	0		
4	MobileNet	0.7125	0.7649	0.6961	0.6826	5	0		
6	MobileNetV3Large	0.4250	1.5067	0.3065	0.3684	5	0		
5	MobileNetV3Small	0.2479	1.7233	0.1870	0.2401	5	0		

ubuntu@ip-10-1-3-133:~/Jupiter_d1/Final_project/DEEP_LEARNING_6303_GROUP3/Code/Trashnet/code\$ python3 trashnet_v5_JK.py

My version: need to see if increasing alpha helps with model performance or just an accident, so here are some new changes:

- Define alpha values to test for class weight damping: 0,0.2,0.5,0.8,1
- Shuffle train data, don't shuffle val data
- Outer loop for testing diff alpha values: so that the result comes inside the loop.

```
# --- Outer loop for testing different alpha values --- NEW JK
for alpha in alpha_values_to_test:
    print(f"\n\n{'*20'} TESTING ALPHA = {alpha:.2f} {'*20'}\n")

    class_weight_damped = 1.0 + alpha * (base_weights - 1.0)
    class_weight_dict = dict(enumerate(class_weight_damped))
    print(f"Using Class Weights for alpha={alpha:.2f}: {class_weight_dict}")
```

- Usually, using alpha level can help with enhancing model performance, because alpha adjusts the balance between classes, similar to weighted cross-entropy but with added flexibility.
 - o Highly skewed datasets: in datasets where the minority class is extremely rare, setting alpha around 0.25 for the majority class and 0.75 for the minority class can be effective.
- not really worth it because it takes way too long to run the script

- alpha level is not automatically uploaded, selected randomly with not very precise calculation or based off results from the model performance. So disregard.
- ///
-
- data augmentation:
 - o adding Sheer_range , brightness_range to ImageDataGenerator used for training and validation
 - to further increase the diversity of the training data by applying more types of random transformations
 - helps model become more robust to variations, reduce overfitting

citation:

Yassin, A. (2024, November 13). Adam vs. AdamW: Understanding Weight Decay and Its Impact on Model Performance. *Medium*. <https://yassin01.medium.com/adam-vs-adamw-understanding-weight-decay-and-its-impact-on-model-performance-b7414f0af8a1>

Code percentage:

Roughly: original script 136 lines

Me editing/modifying/ 180 lines

This is very hard to say because the baseline code is a easier model so since v2 I didn't trace every single line, but I did generated the codes and modified the codes after understanding them. There's no reason for me to lie and say I hard code every single line but I do understand them and with the help of how to make the code writes smoother and better. The architecture is indeed very interesting and it is hard to combine what we learnt in class to a real life case. I enjoy the class.

